

Identifying Corporate Donation Strategies: HTE Estimation Through Causal Forests

true true

2021-05-10

Abstract

The undue influence of political donations on parliamentary voting behavior is well-researched, well-known, and a truism of public perception of American politics. However the reverse – the effect of voting behavior on the direction of campaign funds by major corporate donors – does not appear in prevailing literature. While the idea that legislators are beholden to their donors interests is colloquial, current research fails to investigate the mechanisms through which donors respond to votes. Therefore, the following project aims to build a model that can predict changes in the level of campaign donations to legislators in the US House of Representatives as a function of their voting behavior in the previous legislative session. The project applies causal forest models to recover the effect of party line voting. The model aims to test the theory that donors invest in political campaigns primarily out of expectations of legislators following party positions. Deviations from this expectation - voting against the party line - is therefore expected to be penalized by donors. In order to help elucidate how machine learning techniques can be employed in causal inference settings, we will highlight the limitations standing in the way of recovering causal effects in the relationship in question and how recent developments in machine learning methods of causal questions can be implemented in observational settings.¹

Introduction

In the wake of the landmark 2010 US Supreme Court decision *Citizens United vs. the US Federal Election Commission*, the influence of corporate money in US politics has risen impressively. In parallel with a tidal wave of money washing over the electoral process is the rising tide of social science literature aiming to capture the effect of money on the voting behavior of legislators. Building off of decades of research into the determinants of congressional voting, recent literature has employed state of the art ML methods to attain predictions with considerable success.

Authors have been able to accurately predict past voting behavior (Smith et al. 2012) and future behavior of non-incumbents (Bonica 2018) using campaign donation history and supervised learning methods. These publications have resulted in more accurate predictions of voting behavior than previous statistical methods, specifically those that draw from spacial models inferring ideological positions of legislators from past voting behavior (Clinton, Jackman, and Rivers 2004)(Poole, n.d.). This has been particularly helpful in providing predictions for non-incumbents without a voting history.

In a study published by Adam Bonica in 2018, random forest and support-vector regressions were used to map predictions of non-incumbents on inferred ideological positions from prevailing spatial models through campaign donations (Bonica 2018). While the models provided vote predictions for this cohort (non-incumbents) just as or more accurate than roll call predictions of legislators with a voting record, he writes, "I exclude corporate and trade PACs from the feature set due to their tendency to mix ideological and strategic motives." The success of ML methods in using campaign donation data indicates that political contributions

¹Link to our GitHub repository

are a powerful predictor voting behavior, however the prevailing literature neglects to address the potential of reverse causality: that voting behavior can have an affect of the direction of campaign contributions. We therefore aim to shed light on this comment made by Bonica, namely the potential ideological or strategic motives of major campaign contributors and how they respond to congressional voting behavior.

Despite a paucity of scientific literature on this issue, current journalistic coverage of US politics demonstrates how common knowledge it is that donors actively respond to legislative outcomes. In February of 2021, Thompson Reuters reports, "Ten U.S. corporations slashed donations to candidates seeking federal office by more than 90 per cent in January, after pledging to cut off giving to the Republicans who supported former President Donald Trump's attempt to overturn his election defeat." (Reuters 2021) Additionally, as a result of anti-easy-voting laws in Georgia, many corporate donors have publicly chided legislative officials and threatened to pull donations from key legislators. (Schwartz 2021)

Whether threatening lawmakers that supported a ham-handed coup in 2021 or those that too aggressively seek to regulate certain industries, donors are responsive actors in the US political system. Implicit in studies on the voting outcomes as function of political contributions is that legislators are somewhat beholden to their donor's interests. Logically, the threat of reducing or canceling contributions must be credible, otherwise legislators would not be incentivised to follow donor interests. And if donors were totally unresponsive to 'disloyal votes' by legislators, research examining the influence of money in parliamentary bodies might itself be moot. We therefore believe that to capture a full picture of corporate money in the electoral process, further literature must reverse its typical methodologies and situate the donor at the center of analysis.

To describe the potential relationship more concretely, legislators may signal to untapped donors their position on new issues, or break with existing donors on important pieces of legislation, running afoul of their contributor's interests. Such action would theoretically have an effect on how donors choose to funnel campaign funds in the future, and over time may lead to substantial shifts in corporate giving as parties stake out new positions over time. As the US party system has experienced a seismic shift over the last 10 years in relation to positions on health care, taxation, finance, and resource extraction, the potential to observe shifts in corporate giving is strong.

Drawing on data of corporate donations to individual candidates (including from the organization itself, family of owners/leading representatives, and associated PACs, or, Political Action Committees), we reverse the focus of prevailing literature to examine how corporate donors respond to voting behavior in the US House of Representatives. Recovering corporate donation strategies and further shedding light the risk-reward structure that legislators navigate a central motivation of this project.

To identify any potential relationship between our independent variable of interest and the resultant change in campaign contributions in response to those votes, we follow a causal identification approach. Casual inference goes beyond transitional statistical or machine learning methods by positing and defending a set of logical assumptions that can affect the design of the model and, if the assumptions are defended credibly, allows the researcher to make causal claims about the relationship in question. Congress members may face tough choices between the opinion of their district, party, or conscience and the needs of their donors, and if they choose to respond to pressures from one stakeholder other than their donor, we believe these votes may cause reduced or canceled funds in subsequent cycles.

Using available data on political contributions made by donor organizations to individual House Members from 2012 to 2018, this project analyzes the effect roll call votes grouped by legislative topic and coded as a percentage of votes followed the party line. These party-line vote ratios are averaged into a single percentage of party voting in that year. We theorize that donors are likely to donate to members based on their party's position on major topics that affect their businesses, and that the probability of following the party line on a given topic is a major signal to donors that will be either rewarded or penalized depending on the legislator's party and the donor in question. Although parties are not completely homogeneous, the rate of party-line votes has been steadily increasing over the last decades, which reasonably indicates that donors may expect parties to deliver more consistent support on various legislative topics.

This also means that breaking with party lines could potentially be a stronger treatment, as is it much less common. The outcome of the model is the difference in donations from one year to the next for a specific legislator and specific donor, with the preceding years' party line vote metric used as the treatment.

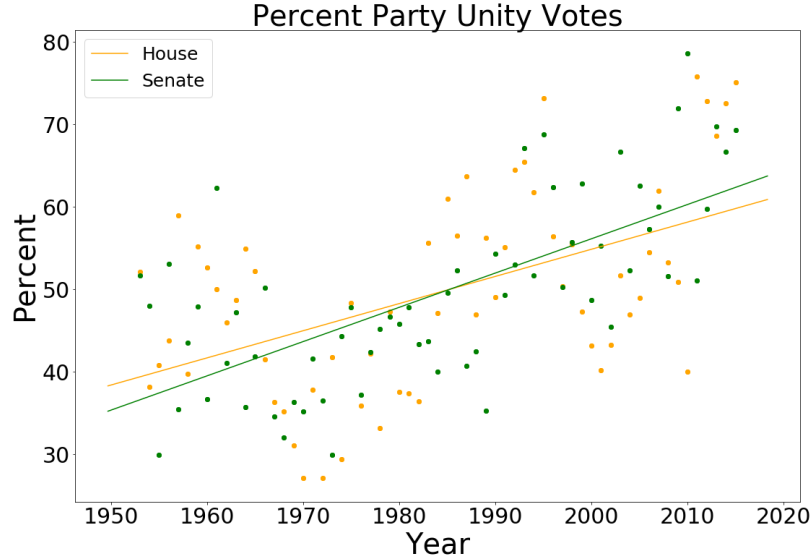


Figure 1: Number of partly line votes in US Congress

Our model includes additional covariates that may moderate outcomes. These covariates are where we theoretically would be able to interpret donor strategies from the model. For instance, we include party identification, which would allow us to interpret the donor's behavior toward each party and we include year dummies to indicate shifts in donor strategies over time. After running our models, we were confronted with major deficiencies in our data collection process that prevents us from making substantial interpretations of our results. Above all, a lack of variability in our outcome variable - the change in donations from one year to the next - severely limited the number of observations in which treatment effects could be observed. Donors proved to be somewhat stable in their donations, and as a result our model outcomes had extremely high confidence intervals. Despite having no meaningful results to interpret, the following project demonstrates how Casual Forest methods can be implemented and how they differ from random forest estimation.

Related Work

The integration of machine learning and causal inference has led to a recent boom within the field of econometrics. As our task seeks to discover potential donor strategies that differ based on metrics such as party affiliation and the specific election cycle, we draw from literature investigating heterogeneous treatment effects within observational data. The relevant literature sketching out possible methodologies has been primarily published within the last 5 years, and as such provides overview of a growing field and the associated ML tools created for such models.

Athey and Imbens 2016 study, "Recursive partitioning for heterogeneous causal effects," serves as a starting point for the balloon in literature on adapting random forest algorithms for causal questions (Athey and Imbens 2016). The authors demonstrated how random forest algorithms can be modified to optimise a treatment effect of subgroups rather than optimise the prediction of an outcome. They demonstrate that these causal effects can be recovered in both trial and observational studies. The main contribution of this study is that using their adjusted methodologies can produce confidence intervals for subgroups through the partitioning of the tree in one sample and the estimation of treatment affects using another. This methodology is called the "honest" approach and will be implemented in our models.

Davis and Heller (2017) apply a similar algorithm to identify heterogeneous treatment effects of youth employment programs on violent-crime arrests and employment. Using data from randomized controlled

trials of a summer jobs program, they find that the applied method is able to identify treatment heterogeneity that a conventional model with interactions effects would have missed. However, running a in-sample, out-of-sample and an adjusted-in-sample model, the authors conclude that the success of out-of-sample models is quite sensitive to the sample size, and causal forests might be better suited for setting with more observations (Davis and Heller 2017).

Wager and Athey (2018) further assessed the properties of causal forests algorithms within a quasi-observational setting. They use a dataset modeled after the National Study of Learning Mindsets, which evaluated the impact of a mindset intervention (instilling a “growth mindset”) on students scholastic achievements, to examine practical and conceptual challenges of causal forests. They find that especially clustered data (at the school level) needs to be considered in casual forest approaches and may lead to substantial bias if not controlled for. They present a set controls that can be implemented in causal forest estimation packages within R (Wager and Athey 2018). A year later, Athey again expanded use case of causal forest algorithms to including IV based approaches and conditional average partial effects estimated within observational and experimental settings.

In April of 2021, David Jacob published a study entitled, "CATE meets ML," which provides a comprehensive overview to various ML techniques that can recover heterogeneous treatment effects, or CATEs (conditional average treatment effects) (Jacob 2021). Together, these papers helped us to understand how these novel methods work and how they can be applied to our causal question: Do donors reward or sanction legislators for voting for or against party lines?

Proposed Method

The collection of our data and construction of our model was guided by a causal approach, specifically using the potential outcomes framework (POF). The POF is a convenient way of demonstrating the fundamental problem of causal inference. In order to be able makes claims on causality, rather than simply correlation, we need to know two things, the outcome for an observation in the control state and the outcome for that very same observation in the treatment state.

$$TreatmentEffect = \mathbb{E}[Y(1) - Y(0)],$$

In the real world, however, only one state can be observed. Either a patient receives a vaccine in a clinical trial or a placebo - but not both. Without being able to observe both states of the world for the same observation, we must compare treatment and control groups who are necessarily different from one another in many ways. Fundamentally, the difficulty of causal inference is either removing methodologically any factor that may influence your outcome other than the treatment or assume away potential confounders. If the resultant methodology and set of assumptions are convincing, causal effects may be recoverable. In randomized controlled trials, this is much easier to manage, however in observational settings where many confounding variables are unobservable, recovering causal effects is a tall order.

In order to derive a causal estimation of the effect of party line voting, we began this project by conceptualizing the causal pathways of the relationship we were hoping to measure. Our data captures the treatment as a percentage of votes along party lines for the individual legislator and our outcome as the change in donations from the previous cycle. We collected these data points over 4 time periods, allowing us to potentially remove considerable bias out of the model.

The graph in Figure(2) is what is called a Directed Acyclic Graph (DAG), which visually represents how variables in our model may be related.

Drawing from seminal work by Imai and Kim (2019), we initially tried to create a model that would satisfy the following four assumptions necessary to estimate causal inference within panel data. (Imai and Kim 2019).

1. **No unobserved time and unit varying confounder:** In figure two, we can see the unobserved confounders at the unit and time level drop out of the model (highlighted in gray) due to the panel

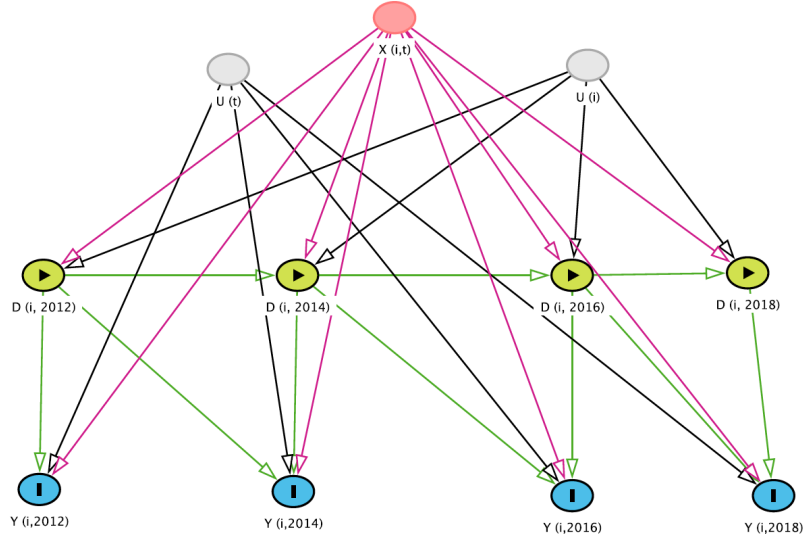


Figure 2: Directed Acyclic Graph (DAG) of causal pathways between votes and donations for the individual legislator

data structure. Having multiple units in single time periods and multiple time periods for single units controls for many potential confounders local to the unit or time period. More problematic for our model would be confounders that vary at the *unit and time level* and might affect donations in the subsequent cycle. This could be changes in committee assignments of a legislator into topics relevant to the donor’s business or gaining a position in the party leadership.

2. **Past outcomes do not influence current ones:** The only way that this would effect our model is if legislators did not spend all of their money from their campaign and carry it over to the next. While this is often true, our data is not taken from the legislator, but rather from the donor, meaning that it includes only the funds disbursed each year.
3. **Past outcomes do not affect current treatment:** This is the most difficult assumption to hold within our data structure, and generally requires the use of an instrumental variable that can reasonably predict the treatment and has an affect on the outcome exclusively through that treatment. Unfortunately without an IV that can predict our party-line vote metric while remaining otherwise independent from the outcome, we could not have made exclusive claims on causal effect. Practically, there will always be a chance that the party-line percentages are determined by donations in the previous year, rather than the other way around.
4. **Past treatments do not affect current outcomes:** While past treatments do influence future outcomes (and is in fact the main effect we want to measure), we would only have to compensate by measuring the effect of 2012 party line votes as a predictor of 2014 donations. This is already incorporated into the the DAG and into the data structure.

While the model structure in the DAG proves sound for traditional statistical methods, we ran into severe limitations of our software in working with panel data, and as a result were forced to alter our model strategy from looking at individual treatment effects (at the level of the legislator) to conditional treatments effects, which are ‘conditioned’ by some moderating variable that splits the observations into subgroups.

To do this, we implemented the CausalForestDML method within the EconML package developed by Microsoft Research. It is the first software package in Python that provides a framework for estimating heterogeneous treatment effects (also known as conditional average treatment effects) for observational data. Models within the package are geared toward providing subgroup treatment effect estimates, which meant that our goal of understanding donor behavior in relation to certain groups, for instance Democrats or Republicans, could be achieved.

For our model, we chose to implement a Causal Forest Regression, which has two major differences to random forest models. Firstly, while a typical regression tree splits the data on a variable and a value that reduces a measure of uncertainty (i.e. MSE), a causal tree will make a split on a variable and a value that maximizes the difference in the outcome across the treatment and control conditions within a leaf. While a random forest makes prediction, the causal forest calculates the treatment effect for a subgroup at each split. The second major difference regards how the data is split. An "honest" causal tree will split the data into two groups, a splitting sample and an estimating sample. The process above, of calculating treatment effects, will be done within the splitting sample until the tree has been constructed. This tree is then applied to the estimating subsample, which lets the observations drop down the tree until they reach a terminal node. Once the data is in place, the CausalForestDML method will implement an estimation function to recover treatment effects within each leaf. The result is an "honest" causal forest, where the resultant point estimates are normally distributed.

In terms of our potential outcomes framework, our model no longer attempts to recover individual average treatment effects, but rather conditional average treatment effects (CATE). The "conditional" element is subgroup in which treatment effects are calculated. The estimate can be defined as:

$$\tau(x) = \mathbb{E}[Y(1) - Y(0)|X = x],$$

Where X could be any moderator of the treatment variable, such as gender, party affiliation, or ethnicity. To implement this model within our data, we followed case examples provided by the github repository for the EconML package (Microsoft 2021) and from (Naushan 2021). In order to assist in interpreting our models, we additionally implemented the SHAP (SHapley Additive exPlanations) package in Python to uncover potential drivers of our treatment effects.

Experiments

Data: Our analysis relies on data coming from three different sources. First, we include data on all Members of the House of Representatives (MoHs) who served in all sessions along our timeline, which encompasses the 113th to the 116th Congress. In total this is 224 house members. We made this decision to restrict our observations when we were following a panel data estimation strategy. To get complete lists of all people who served in the respective sessions in the observed time frame, we are using the comprehensive data sets of the Comparative Legislators Database provided by Munzert and GÅ¶bel (Göbel and Munzert 2021), which we adjusted for our needs ("01_Creating_MoH_Dataset.ipynb"). In addition to a complete list of MoH's names, we use the database primarily for information regarding the gender, race and party affiliation of MoHs, which are examined regarding their predictive power as features in Random Forest Models and serve as moderators in the Causal Forest Models.

Second, we include data on roll call votes taken from GovTrack.us (Civic Impulse, LLC 2021) that contains the individual votes of all Members of the House (MoHs). The bills that are subject of roll call votes are sorted into categories which allows us to examine in a more targeted manner in which area individual voting behavior by MoHs has an effect on donor behavior. For our main independent variable we aggregated the results of all roll call votes in several categories (given by GovTrack) that were chosen based on the business model of the respective donors examined in our analysis. Based on the roll call votes in each category a variable is created that indicates the ratio of votes along the party line for the individual legislator. The party line is defined as the vote of the majority of MoHs of the same party on the respective bill ("02_Cleaning_Vote_Data.ipynb"). By having the ratio of votes along the party line for each MoH in several different categories, we are able to analyze in which areas donors respond to individual votes and in which areas diverging voting behavior

has little impact on future donations. As we are looking at party lines, our analysis does not include votes by independent candidates.

Third, the donation data is taken from OpenSecrets.org which is run by the Center for Responsive Politics (The Center for Responsive Politics 2021). It categorizes donations per election cycle, meaning that donations for 2012 will be to candidates serving in the congressional session beginning in 2013. This structure of the data is why we will be lagging out the predictor variable, so that the voting percentages in session beginning 2013 will be used to predict values for donations in the 2014 election cycle and so on. For our dependent variable we are taking the change in donations from a certain donor to an individual MoH between two sessions. (03_Merge_Donations_MoH(Walt).ipynb).

In all, there are 30 categories of legislation available on GovTrack, however we were unable to automate the scraping of the data and as a result were only able to include 5 legislative categories in our data set. As will be discussed in subsequent section, the lack of data meant that we could not reliably run tests to determine whether donors were specific to a subset of a broad range of categories, nor could we know whether the party line vote metrics for the 5 categories (mostly financial related.) are not systematically different from overall voting behavior. Similarly, we were not able to automate the data collection process for the donor data either. As a result we only have data for two donors, and neither of them exhibit high variability in our outcome variable.

Our two donors are Walt Disney Inc. and Bloomberg LP, and our categories are Commodities Markets, Financial Crisis Stabilization, Insurance Industry and Regulation, Real Estate Business and Television and Film. Both donors are plausibly related to at least two of these categories, and without automating the data scraping process we could not expand our data any further.

Software: We ran our models on Python 3.8.5, and used JupyterLab 2.2.6 as interface to write code and conduct our analyses. Moreover, we used GitHub to collaborate on the project and for version control. EconML and SHAP were used to run the main causal models and provide interpretable output.

Random Forest: In order to better investigate the relationship between the voting behaviour of MoHs in certain legislative categories and changes in donation patterns, we ran several basic Random Forest Models as initial step of data exploration. First, we examine patterns in the voting behaviour of Walt Disney based on the legislative category "Television and Film", and the four other categories that can be attributed to the financial sector (07_Random_Forest_Disney.ipynb). With the outcome variable encoded into a discrete categorical variable that takes on 0 in case of an decrease of donation, 1 in case of no change in donation and 2 in case of an increase in donation between two years, we run a RandomForestClassifier that yields an accuracy of 0.6697 in a baseline model with `n_estimators = 1000` and `max_leaf_nodes = 10`. The accuracy can be increased up to 0.7030 by running another model with tuned hyperparameters using RandomizedSearchCV (Chen 2021) with `n_iter = 100`, `cv = 3`.

A RandomForestRegressor that takes the outcome as a discrete variable in absolute terms, i.e. changes in donations by Walt Disney between two years, returns a RMSE of 2763.61 in a baseline model with `n_estimators = 500`. Again by using RandomizedSearchCV with `n_iter = 100`, `cv = 3`, the RMSE can be brought down to 1472.34 for the Walt Disney donation data (MAE and MSE show a comparable reduction in error).

Figure(3) exemplifies the grid search cross-validation process on a two-dimensional level by visualizing potential combinations of `max_depth` and `max_features` (in this example via GridSearchCV instead of RandomizedSearchCV).

Figure(4) illustrates the importance of all features used in the Random Forest Regression Model trained on the Walt Disney donation data. The underlying values are based on the best random model determined by RandomizedSearchCV.

As expected, the variable encoding the ratio of voting along parties lines in the category "Television and Film" proves to be most important (21.22%) for predicting changes in Walt Disney's donation behaviour.




Figure 3: Random Forest Regressor cross-validation via GridSearchCV (Koehearsen 2018), CV score = -0.049983, at `max_depth = 1` and `max_features = 3`




Figure 4: Feature importance for Walt Disney donation data based on best random model via RandomizedSearchCV (Chen 2021) (with `n_iter = 100`, `cv = 3`)

Voting along party lines in the other four examined categories, however, also entail high predictive power with values ranging from 6.98% up to 16.30%. Unsurprisingly, both party indicators prove to be more important than other MoH-specific characteristics as gender or ethnicity. However, all MoH characteristics are significantly less important compared to most of the legislative categories. Moreover, there are noticeable differences between the indicators for different years that are included in the model, with the relevance of a certain voting behaviour in 2016 being almost twice as important as in 2012.

Secondly, we examine whether a similar clear relationship between Bloomberg LP’s donation patterns and the voting behaviour of MoH in legislative categories associated with the financial sector prevails (07_Random_Forest_Bloomberg.ipynb). Again encoding the outcome as a discrete categorical variable, a baseline RandomForestClassifier with `n_estimators = 1000` and `max_leaf_nodes = 10` yields an accuracy of 0.9598 which is notable higher compared to the results for Walt Disney. Running a second model with tuned hyperparameters using RandomizedSearchCV with `n_iter = 100`, `cv = 3` does not further increase accuracy.




Figure 5: Feature importance for Bloomberg LP donation data based on best random model via RandomizedSearchCV (Chen 2021) (with `n_iter = 100`, `cv = 3`)

Figure(5) illustrates the importance of all features used in the Random Forest Regression Model trained on the Bloomberg LP donation data. The underlying values are based on the best random model determined by RandomizedSearchCV.

Again, the ratios of voting along party line in different legislative categories prove to be the most important features regarding the prediction of changes in Bloomberg’s donation patterns.

Causal Forest: The previous models indicated that party line voting was at least a reasonable predictor of the data we had. With this in mind, we moved forward to our causal models to recover treatment effects along the categorical variables we had in the data.

As the analysis of features importance for both Disney’s and Bloomberg’s donation data has shown, legislative categories that are related to the respective business segment have a higher predictive power compared to other categories and compared to MoH-level characteristics. As we had decided to code our treatment variable as an average of party line vote percentages across the topics we had within the data, the results from the Random Forest models dampened our expectations of the CATE estimations as they would be driven by categorical features with low predictive power.

In running the causal models, we shift from estimating model accuracy to estimating conditional treatment effects of a single party-line vote ratio. Using CausalForestDML, we implemented an ‘honest’ causal forest model with covariate dummies for year, ethnicity, gender and party, and the outcome variable as the difference in donations between years. We chose a MultiTaskLassoCV estimator for both fitting the outcome (model_y) and the treatment to the features (model_t) as both our treatment and our outcome variables are continuous (MicrosoftResearch 2019).

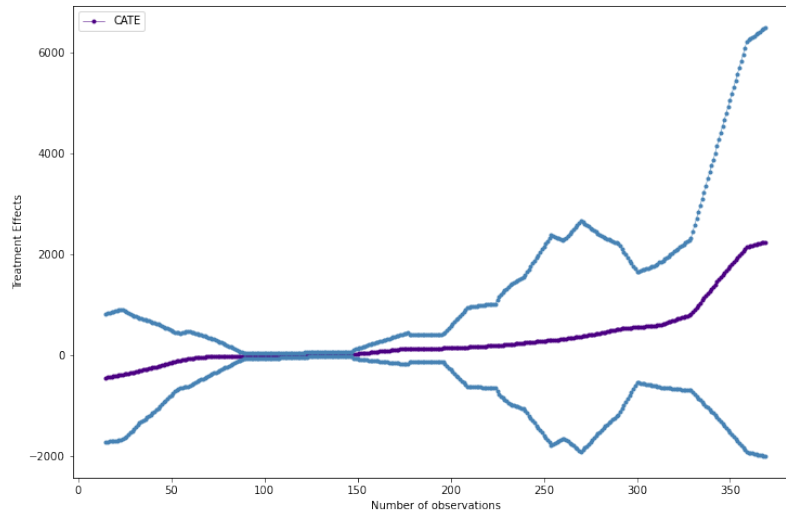


Figure 6: Conditional average treatment effects and confidence intervals for Bloomberg LP donation data

As Figure(6) illustrates, the confidence intervals of the conditional average treatment effects on changes in Bloomberg’s donation patterns vary widely, with a large section in the middle hugging no effect. This is primarily due to the sheer lack of variability in the outcome variable. This also potentially explains why the Bloomberg Random Forest model was able to predict so well: a vast majority of the observations showed no change, with only a few outliers.

Figure(7) displays the SHAP values, i.e. the impact of covariates on the outcome variable, based on Bloomberg donation data. To be precise, the figure illustrates three things: First, feature importance is ranked in descending order, meaning that "Year_2014" is the top feature in the model. The color indicates high or low values of that feature and the direction indicates whether the values contributed positively or negatively to the outcome. For example, the covariate "party_short_D" is a dummy variable that takes on values 0, i.e. a low feature value (blue), and 1, i.e. a high feature value (red). Looking at the SHAP value, we can see that red points, i.e. being a member of the Democratic party, have an positive impact on the models output, i.e. a positive change in donation money from Bloomberg, while blue points, i.e. not being a member of the Democratic party, have an negative impact on the outcome variable. This aligns with the fact that Bloomberg LP is a highly partisan donor that mostly gives to Democrats.²

We additionally ran a causal forest model on the Disney donation data. Figure(8) displays the respective confidence intervals around the CATEs. As we can see the confidence intervals do vary less over the number of observations but do include zero at all times. This is likely a result of greater outcome variability within the Walt Disney data. Figure(9) shows us the SHAP values for the features for the causal Disney model.

²See allocation of Bloomberg LP’s political contributions




Figure 7: SHAP values (impact on model output) based on Bloomberg LP donation data




Figure 8: Conditional average treatment effects and confidence intervals for Walt Disney donation data




Figure 9: SHAP values (impact on model output) based on Walt Disney donation data

Analysis

The unfortunate reality is that the data the models is so sparse that any attempt to recover heterogeneous effects is impossible. Moreover, the interpretation of certain results are meaningless given that outcomes are likely to be generated by randomness in the data than any real causal relationship.

Firstly, while our initial experiment in testing the feature importances of legislative categories against specific donors was supposed to indicate whether we had matched a donor to the categories that we were able to manually scrape, it is extremely likely that any category could simply increase the prediction of the random forest model by chance rather than any association. For instance, the Bloomberg donation outcomes were strongly predicted by the voting metric in the Film and Television category. This could have been by chance or because the legislator had similar voting habits across all categories - in effect they either a consistent partisan or non-partisan donor. Without all the legislative categories, it was our choice then to simply average the party-line vote percentages into one single treatment variable.

Perhaps the most problematic of our data was the lack of variation in the outcome. Our outcome variable was only non-zero for roughly 10 per cent of our data in the Disney data set, and even less for Bloomberg. The result is that our model is driven by variability in only a few observations. The data bottlenecks we encountered in scraping the data unfortunately prevented and solid basis from which to run the models, and as such our results are somewhat meaningless.

Additionally, for what observations that did have changes in the outcome, we did not have solid covariates to run against them. For instance, there are many use case examples of EconML that have only 100 to 200 observations, however have a large number of covariates to run against them - in fact this is what the EconML is particularly suited for. In that sense, the Disney data could have theoretically provided enough variability in the outcome for treatment effects to be estimated given that we additionally had strong covariates. The results from our Random Forest models indicated that our cateogrial features did not add significantly to the accuracy of the models and therefore we were not surprised to see such inaccurate CATE estimates among subgroups. Examples of additional features we had hoped to include were: relevant committee assignment, party leadership position, competitiveness of district, and margin of victory in primary. All of these are potential observable confounders that we were not able to successfully integrate within our data sets, and

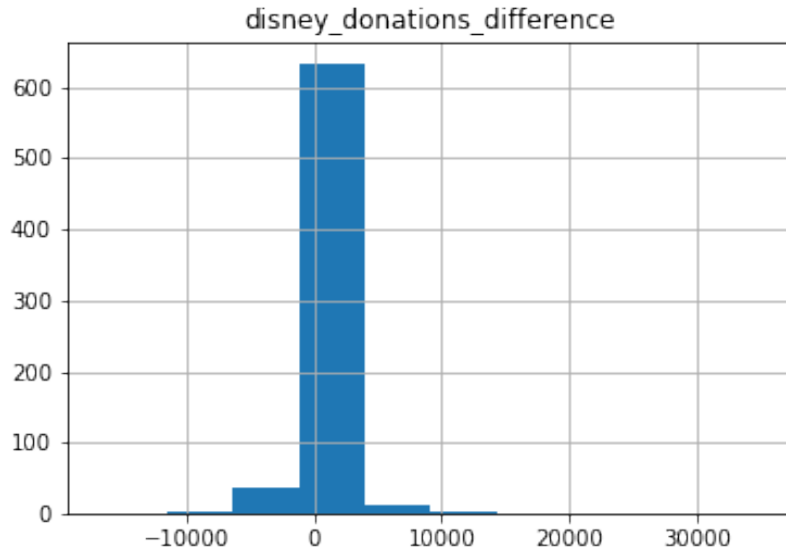


Figure 10: Distribution of Outcome Values

would logically have a larger effect on donor motivations than either ethnicity or gender. While the year features could have potentially revealed changes in treatments effects across important election cycles, we are hesitant to make substantial claims given the poor quality of the data.

Conclusions

Because of the lack of data, our study unfortunately has no key findings to present. While this is extremely disappointing given the interesting subject matter and the exciting new field of casual ML methods, we have walked away from the project with a deeper understanding of how ML methods can be adapted to support causal inference approaches. Moreover, we are walking away with a solid understanding on how the CausalForestDML method works and how causal forests differ from random forest estimation. The primary limitation of our work is, as we have repeated so often, the fundamental lack of usable data. The EconML package is specifically designed to handle use cases with not only a high number of observation but also a high number of covariates. Given that we are both beginners to machine learning, taking observational data already collected and cleaned would have given us the space to delve further into multiple methods within the package and explore model tuning. Additionally, having a model with solid covariates would have opened space to delve into model explainability, which is a key development for the integration of machine learning methods into questions of causal inference. Our focus on creating a convincing causal approach—and collecting and cleaning the data to fit this approach—reflected our experience with causal inference within statistical courses. In hindsight, letting go of our focus on finding a model suitable to panel data would have sent us in a much more productive direction. Our key lesson from these failures is that more data provides you more flexibility, and would have prevented us from running what is actually demonstration of a causal forest project workflow rather than recovery of casual effects. An avenue for future work would be to expand our data to include all thirty legislative categories as coded in our GovTrack source as well as additional covariates that we have mentioned. Resolving the issue of panel data within the EconML package is also an avenue for future work. So far as we know, there is no framework for estimating the treatment effects within the package that accommodates such a data structure. Pulling data from many donors and creating legislator level changes in donations across a specific subset donors could have increased outcome variability among and allowed us to make generalizations for donor behavior across the board. With these issues overcome, we might have ended up with the project we'd hoped for.

Acknowledgements

Without the many blog posts published in the last few months on the subject of causal inference and machine learning, we would have been completely lost. We would additionally like to thank our Slava, Hannah, and Huy for guiding the course.

- Athey, Susan, and Guido Imbens. 2016. “Recursive Partitioning for Heterogeneous Causal Effects.” *Proceedings of the National Academy of Sciences* 113 (27): 7353–60.
- Bonica, Adam. 2018. “Inferring Roll-Call Scores from Campaign Contributions Using Supervised Machine Learning.” *American Journal of Political Science* 62 (4): 830–48.
- Chen, James Ming. 2021. “An Introduction to Machine Learning for Panel Data.” *International Advances in Economic Research*, 1–16.
- Civic Impulse, LLC. 2021. “GovTrack Database.”
- Clinton, Joshua, Simon Jackman, and Douglas Rivers. 2004. “The Statistical Analysis of Roll Call Data.” *American Political Science Review*, 355–70.
- Davis, Jonathan, and Sara B Heller. 2017. “Using Causal Forests to Predict Treatment Heterogeneity: An Application to Summer Jobs.” *American Economic Review* 107 (5): 546–50.
- Göbel, Sascha, and Simon Munzert. 2021. “The Comparative Legislators Database.”
- Imai, Kosuke, and In Song Kim. 2019. “When Should We Use Unit Fixed Effects Regression Models for Causal Inference with Longitudinal Data?” *American Journal of Political Science* 63 (2): 467–90.
- Jacob, Daniel. 2021. “CATE Meets ML: Conditional Average Treatment Effect and Machine Learning.” *IRTG 1792 Discussion Paper 2021-005*.
- Koehrsen, Will. 2018. “Hyperparameter Tuning the Random Forest in Python.” Towardsdatascience.com. 2018. <https://towardsdatascience.com/hyperparameter-tuning-the-random-forest-in-python-using-scikit-learn-28d2aa77dd74>.
- Microsoft. 2021. “Automated Learning and Intelligence for Causation and Economics: EconML.” AL-ICE (Automated Learning and Intelligence for Causation and Economics). 2021. <https://github.com/microsoft/EconML>.
- MicrosoftResearch. 2019. “Econml.dml.CausalForestDML.” Econml). 2019. https://econml.azurewebsites.net/_autosummary/econml.dml.CausalForestDML.html.
- Naushan, Haaya. 2021. “Causal Machine Learning for Econometrics: Causal Forests.” Towardsdatascience.com. 2021. <https://towardsdatascience.com/causal-machine-learning-for-econometrics-causal-forests-5ab3aec825a7>.
- Poole, Keith T. n.d. “Howard Rosenthal. 2007.” In *Ideology and Congress*.
- Reuters, Thomson. 2021. “Big u.s. Companies Slash Donations to Politicians After Trump Election Challenge.” Reuters. 2021. <https://www.reuters.com/article/us-usa-politics-corporate-idUSKBN2AL013>.
- Schwartz, Brian. 2021. “Companies Quiet on Whether They Will Keep Donating to GOP Supporters of Georgia Voting Law.” CNBC. 2021. <https://www.cnbc.com/2021/04/01/georgia-voting-law-corporate-donations-to-gop-under-scrutiny.html>.
- Smith, Samuel, Jae Yeon Baek, Zhaoyi Kang, Dawn Song, Laurent El Ghaoui, and Mario Frank. 2012. “Predicting Congressional Votes Based on Campaign Finance Data.” In *2012 11th International Conference on Machine Learning and Applications*, 1:640–45. IEEE.
- The Center for Responsive Politics. 2021. “Open Secrets Database.”
- Wager, Stefan, and Susan Athey. 2018. “Estimation and Inference of Heterogeneous Treatment Effects Using Random Forests.” *Journal of the American Statistical Association* 113 (523): 1228–42.